


CREATIVITY 15% / IMPACT POTENTIAL 20%

TapeDeck

A failed agent run needs
more than a chat log

Save the experiment as an executable image

Local Qwen3.5 2B + Pharo + MongoDB Atlas



Agent.image

Real Qwen3.5 2B fixes the bug

Mode	Calls	Verifier
Record	4	Pass
Replay	0	Not rerun
Live fork	3	Pass
Fresh run	4	Pass

Actual code change

src/sum.js

– return a – b;

+ return a + b;

Independent request counter. All three live verifiers pass.

One fixture, not a benchmark. Replay does not rerun the verifier.

Replay has to preserve state

01 Request interception

Normalize requests and detect divergence

02 Recorded tool outcomes

Replay behavior, not another model answer

03 Git checkpoint recovery

Restore files before the live continuation

Observed fork: 1 replayed step + 3 live model calls

MongoDB is the recovery store

Companion	Size	SHA-256 prefix
Agent.image	70.40 MB	4ed9a1c282
archive.json	4.28 MB	110e2f979e
Agent.changes	0.32 MB	2bda01df7b
Pharo .sources	45.28 MB	4f93dab730

GridFS

.files
.chunks

Snapshot manifest

IDs, sizes, hashes

Publish the manifest last. Verify every download before restoring.

Actual companion metadata from the prior verified Atlas snapshot

The downloaded image reopens

Actual re-verification output

```
IMAGE REOPENED runs=32 files=491 gitBundles=6  
PAYLOADS VERIFIED 491 file hashes + 6 Git bundles  
RESULT PASS
```

4 / 4 companion hashes match the Atlas manifest

Prior Atlas snapshot, rechecked locally. New Qwen image: 4 runs, 76 files, 1 bundle.
The new image is local; its Atlas upload remains blocked.

IMPACT POTENTIAL 20%

Debug the run, not just the log

Agent developers

Recover a failure with its files and checkpoints

Evaluation teams

Compare a deliberate live alternative

Teammates

Reopen the experiment without new replay inference

Bring one failing run.

github.com/gfade/tapedeck_ai